

# IST 718: Big Data Analytics

## UNIT 4-1 Spark Computing

*Note: MATERIAL here is from a mix of Readings & sources as listed in References Slide(s)*

---

# 01 Spark Computing

---

Apache Spark

Spark Components and Functionalities

Spark RDDs

Spark DataFrames

Spark Operations

Spark Mllib

Spark SQL

# Spark and Hadoop

---

Hadoop is a **framework** for distributed processing and is comprised of YARN, MapReduce, HDFS and common modules.

Hadoop **ecosystem** adds to its capabilities: HBase, Storm, Mahout, Neo4j, Graph, etc.

With capabilities in file distribution, data security, resource management and DR, Hadoop is a **de facto standard** for Big Data.

MapReduce processing is **disk-based**, scaling by adding cheap computers and cheap disk.

Largest known cluster: Yahoo 42,000 nodes. Hadoop MapReduce uses **batch processing** and built ideally to crawl through web sites.

Spark is a **general-purpose** data processing engine that can run on its own or in Hadoop clusters through YARN.

Spark **capabilities** include SQL, streaming, machine learning and graph processing.

Spark is not a replacement for Hadoop but can provide an **alternative to Hadoop MapReduce** under several scenarios.

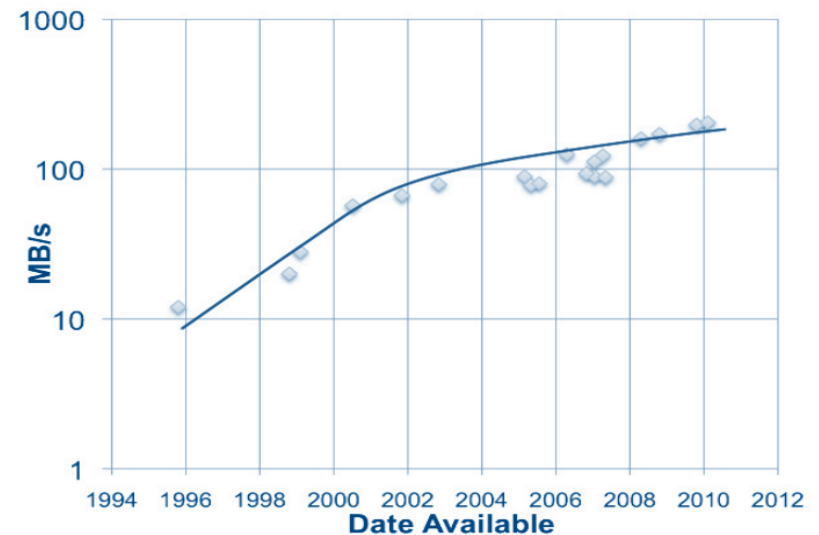
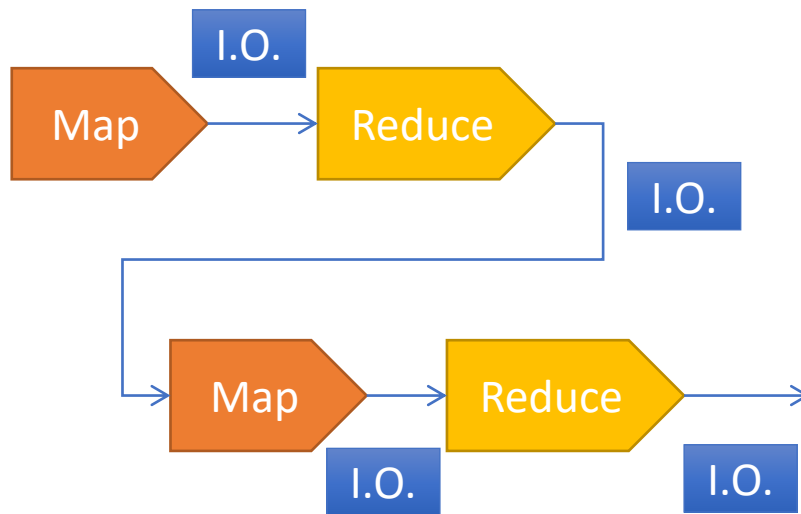
Spark provides rapid **in-memory** processing and scaling comes at a cost (memory).

Largest known cluster: 8,000 nodes. Spark excels at **streaming** workloads, interactive queries, and machine-based learning.

**Note:** Hadoop and Spark can co-exist and is often the recommended solution for new implementations.

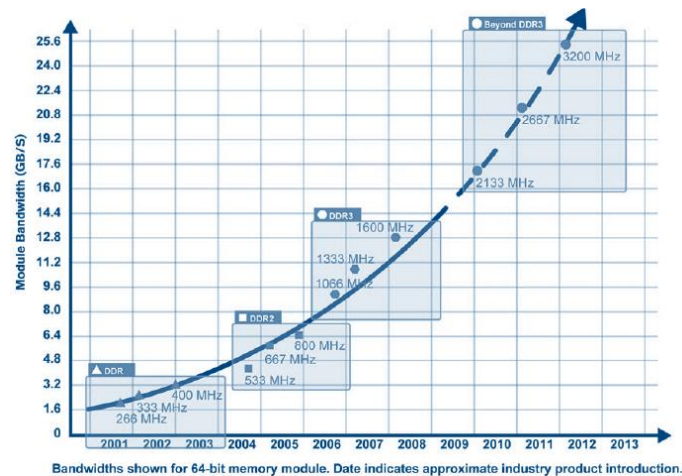
# Hadoop

- Traditionally:
  - Hadoop uses single programming model: MapReduce
  - It only works with data on hard drives



# Spark

- RAM bandwidth has been increasing exponentially
- Spark can perform in-memory computations

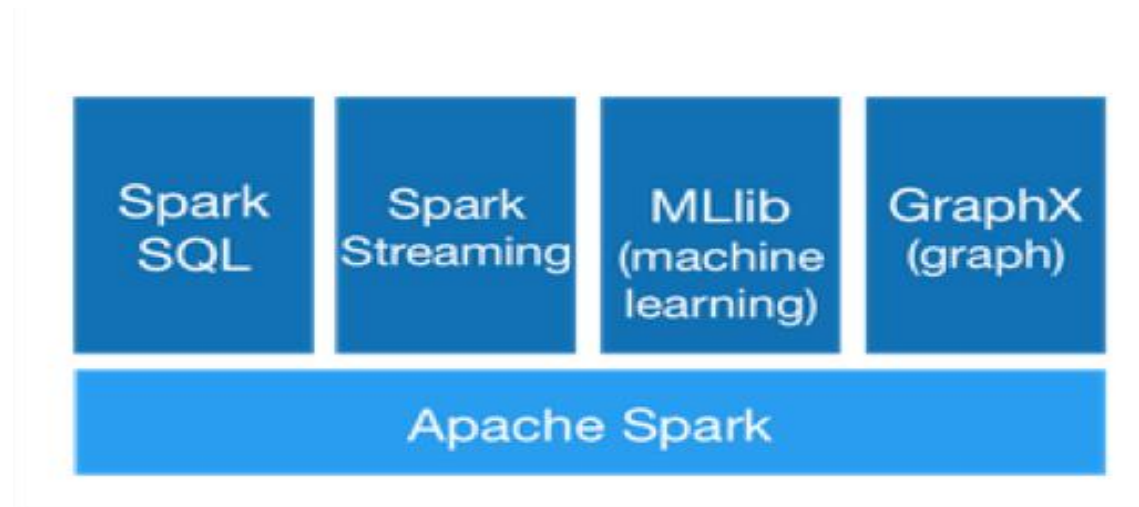


# What is Spark?

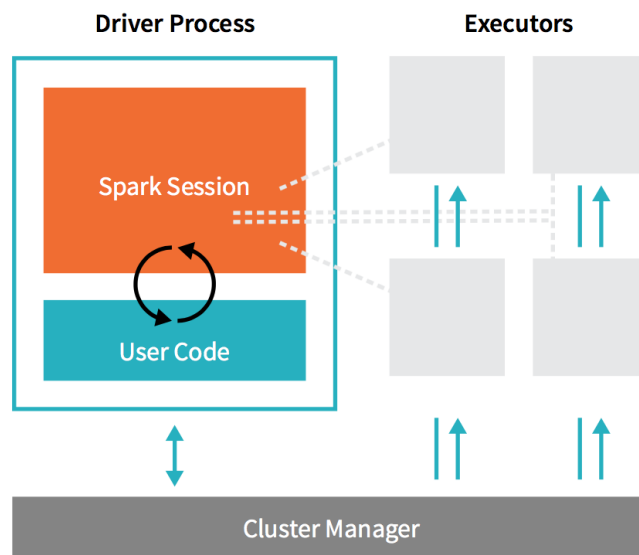
- **Apache Spark** is a fast, in-memory analytics system
- Spark has several high-level tools, including:
  - **ML**: a machine learning library
  - **Spark Streaming**: enables high-throughput, fault-tolerant stream processing of live data streams
  - **Spark SQL**: runs SQL and HiveQL queries
  - **GraphX**: an API for graphs and graph-parallel computation
- Spark can be executed in two ways:
  - Standalone
  - With a cluster manager such as YARN

# Spark

“Apache Spark™ is a fast and general engine for large-scale data processing.”



# Spark Architectural Components



- Spark applications consist of a driver process and a set of executor processes.



# Spark Driver Process

- The driver process runs the application's "main" program
- Runs on one of the nodes in the cluster
- Is responsible for three things:
  - Maintaining information about the Spark applications
  - Responding to a user's program or input
  - Analyzing, distributing, and scheduling work across the executors

# Spark Executor Process

- Responsible for actually carrying out the work that the driver assigns to them.
- Each executor is responsible for two things:
  - Executing code assigned to it by the driver
  - Reporting the state of the computation on that executor back to the driver node

# Cluster Manager

- Controls the physical machines
- Allocates resources to Spark applications
- Three Cluster Managers
  - Standalone
  - YARN
  - Mesos

# Spark 1.6+ vs 2.0+

- Spark 1.6+ relied on transformations on arbitrary datasets known as Resilient Distributed Datasets (RDDs)
- Spark 2.0+ defines more structure in the form of DataFrames, which are similar to tables in SQL and data.frames in R
- The newest versions of Spark define DataSets which are statically-typed DataFrames (more structure)
- With more flexibility comes greater power but less performance
- Future Spark machine learning will work with DataFrames and not RDD.

# Classic Spark v. <2.0

# The SparkContext

- The **SparkContext** object performs the following tasks:
  - It connects to the ResourceManager (like YARN) and asks for resources on the Hadoop cluster,
  - Starts executors on the worker nodes in the cluster that the ResourceManager allocated for the Spark application,
  - Sends the application code to the executors,
  - And finally, it sends tasks for the executors to run
- **SparkContext** is represented by the **sc** object

# Spark RDDs

- An RDD (resilient distributed dataset) is a fault-tolerant collection of elements on which operations can be performed in parallel.
- An RDD is an immutable collection that represents:
  - A dataset...
  - ...broken up into a list of partitions
  - A list of dependencies on other RDDs
  - An optional list of preferred block locations for an HDFS file
  - Read-only

# Important RDD Concepts

- Lineage
  - Information about how an RDD is derived from other datasets or other RDDs
  - RDD is not necessarily materialized all the time due to lazy execution
  - Lineage captured on disk as “lineage graph”
- Persistence
  - Indicate which RDDs which need to keep in memory for reuse
  - User can call “persist” method
- Partitioning
  - RDD elements can be partitioned across machines based on a key in each record



# Sample RDD Lineage Graph

- r20 depends on many other RDDs
- See: <https://jaceklaskowski.gitbooks.io/mastering-apache-spark/spark-rdd-lineage.html>

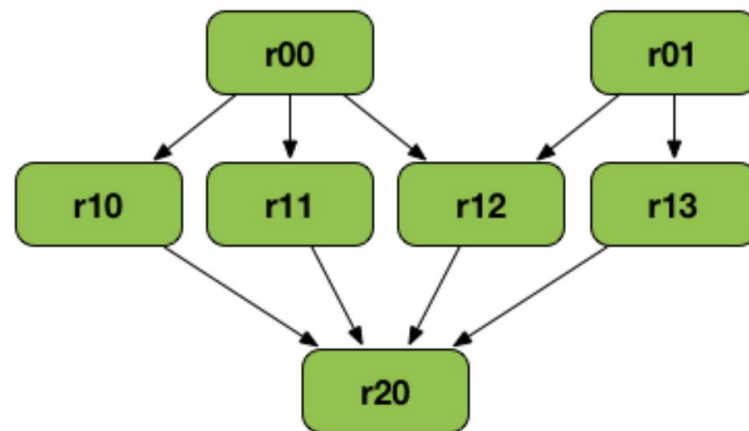


Figure 1. RDD lineage

# Creating RDDs

- Can create an initial RDD by applying a transformation to data on disk
- Can create an initial RDD from a code object
- Example ways to create an RDD in Spark:
  - Use the parallelize operation to convert an existing code object into an RDD
  - Use textFile operation to convert a text file on HDFS into an RDD
  - Use sequenceFile operation to convert a binary file on HDFS into an RDD

# Example: Creating RDDs

```
from pyspark import SparkContext
```

```
myarray = range(1,20)  
myarray
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19]
```

```
dist_array = sc.parallelize(myarray)  
dist_array
```

```
ParallelCollectionRDD[0] at parallelize at PythonRDD.scala:223
```

# RDD Persist

- Spark's persist method

- Indicate which RDDs to reuse
- Indicate if you want to replicate across machines
- Indicate priority of which in-memory data to spill to disk first

- Example

```
logsrdd = sc.textFile("hdfs://user/mark/logdata")
fatals = filter(lambda s: s.startswith('FATAL'), logsrdd)
fatals.cache()
fatals.count()

# Notes: logsrdd NOT loaded into RAM because of lazy evaluation
# fatalities.cache() = rdd.persist(storageLevel.MEMORY) tries to
persist fatalities rdd in memory
```

# RDD Disk and Recover

- RDD can spill to disk
  - Degrade gracefully (to mapreduce performance)
  - Partitions not in use/lesser use and/or low priority spilled first
- Failure Recovery
  - Recovery is facilitated through redundant RDD block copies across cluster computers.
  - An RDD partition that fails is recovered by the Yarn/Spark Driver
  - Executer applies lineage to prior RDD (or original data on disk) to recover the RDD.

# RDD Checkpointing

- Writes an RDD to disk to save the RDD in a specific state
- After checkpointing, future references the RDD don't need to perform upstream transformations in the lineage.
- Similar to caching except that the RDD is stored on disk instead of in memory.
- Example:

```
spark.sparkContext.setCheckpointDir("/some/path/  
for/checkpointing")  
words.checkpoint()
```

# RDD Operations

- There are two types of operations that can be done on RDDs:
  - **Transformations:** create a new dataset/RDD from an existing one
  - **Actions:** Trigger a transformation and instructs spark to compute a result from a series of transformations.
- Transformations
  - lazy - they do not compute their results right away
- Actions
  - Instantiates the RDD
- See <https://spark.apache.org/docs/latest/rdd-programming-guide.html#transformations> for a list of transformations and actions.

# RDD Dependencies and Transformations

- Dependency / transformation used synonymously in Spark the Definitive Guide text book
- Two Main Types of Dependencies / Transformation
  - Narrow – child partition depends on only one parent partition
    - E.g. map, filter, union
  - Wide – multiple child partitions depend on one parent partition
    - E.g. Join and group-by transformations
    - Materialize intermediate calculations on parent for fault recovery



# Narrow Transformation / Dependency

- Each input partition contributes to one output partition
- See Spark the Definitive Guide, Figure 2.4
- Also called pipelining, spark tries to perform narrow operations in memory

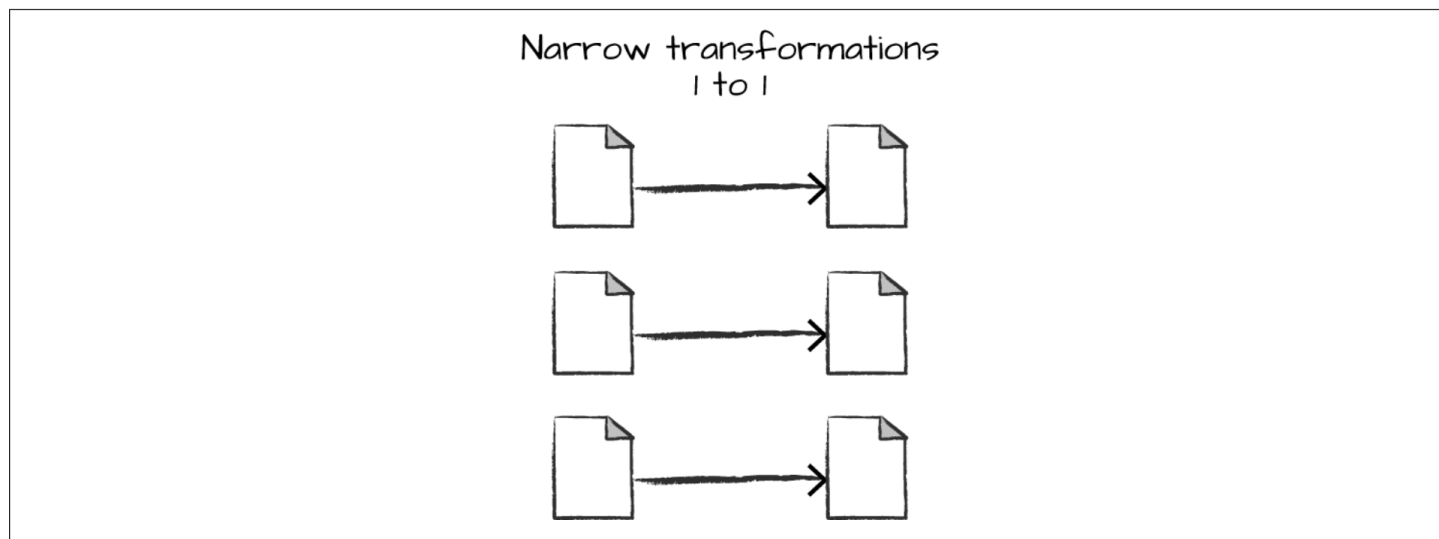


Figure 2-4. A **narrow** dependency

# Wide Transformation / Dependency

- An input partition contributes to many output partitions
- See Spark the Definitive Guide, Figure 2.5
- Also called a shuffle where Spark exchanges partitions across the cluster

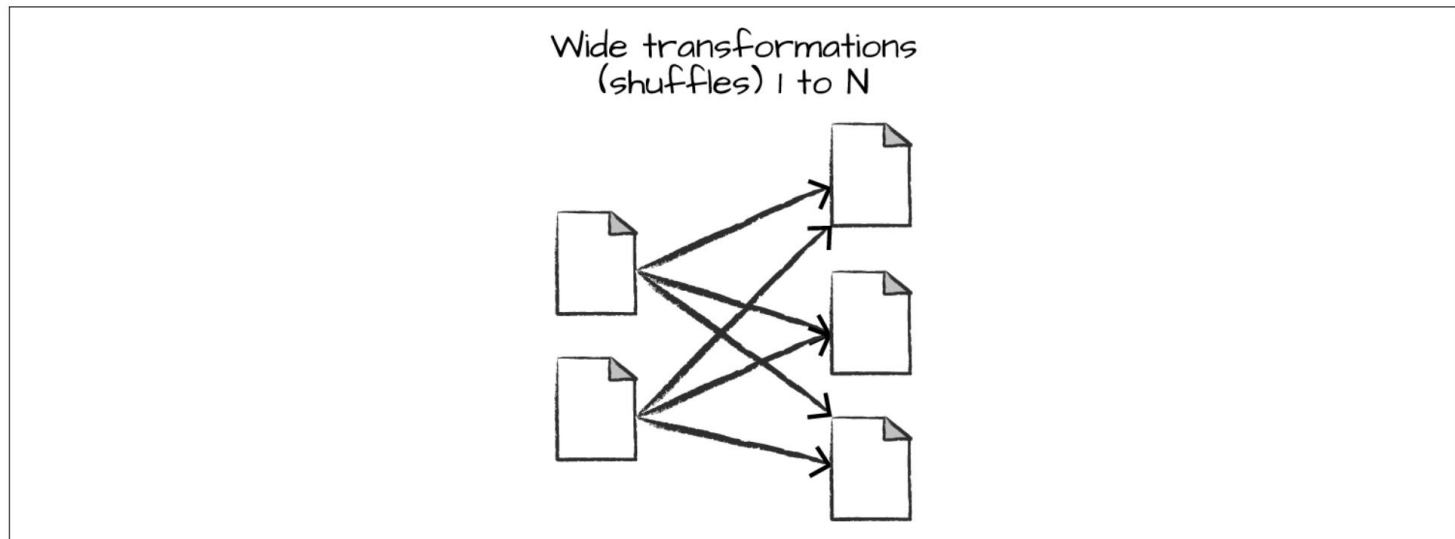


Figure 2-5. A **wide** dependency

# Example Transformations

- **map(func)**: returns a new distributed dataset formed by passing each element of the source through the function *func*.
- **flatMap(func)**: same as **map** but when multiple key-value pairs are returned
- **filter(func)**: return a new dataset formed by selecting those elements of the source on which *func* returns true.
- **reduceByKey(func)**: when called on a dataset of (K, V) pairs, returns a dataset of (K, V) pairs where the values for each key are aggregated using the given reduce function *func*.
- **sortByKey([ascending],[numTasks])**: when called on a dataset of (K, V) pairs where K implements **Ordered**, returns a dataset of (K, V) pairs sorted by keys in ascending or descending order, as specified in the Boolean **ascending** argument.
- **join(otherDataset, [numTasks])**: when called on datasets of type (K, V) and (K, W), returns a dataset of (K, (V, W)) pairs with all pairs of elements for each key.
- **distinct([numTasks])**: returns a new dataset that contains the distinct elements of the source dataset.
- **pipe(command, [envVars])**: pipes each partition of the RDD through the provided shell *command*

# Example: Transformations

```
neg_values = dist_array.map(lambda x : -1 * x)
neg_values.collect()
```

```
[-1,
-2,
-3,
-4,
-5,
-6,
-7,
-8,
-9,
-10,
-11,
-12,
-13,
-14,
-15,
-16,
-17,
-18,
-19]
```

```
large_values = dist_array.filter(lambda y: y > 10)
large_values.collect()
```

```
[11, 12, 13, 14, 15, 16, 17, 18, 19]
```

# Example Actions

- **reduce(*func*)**: Aggregate the elements of the dataset using a function *func* (which takes two arguments and returns one). The function should be commutative and associative so that it can be computed correctly in parallel.
- **foreach(*func*)**: Runs the function *func* on each element in the dataset.
- **count()**: returns the number of elements in the dataset.
- **first()**: returns the first element in the dataset.
- **take(*n*)**: returns an array with the first *n* elements of the dataset.
- **saveAsTextFile(*path*)**: writes the elements in the dataset out to a file in HDFS (or some other file system).
- **saveAsSequenceFile(*path*)**: writes the elements to HDFS in the **SequenceFile** format.

# Example: Actions

```
large_values = dist_array.filter(lambda y: y > 10)  
large_values.collect()
```

```
[11, 12, 13, 14, 15, 16, 17, 18, 19]
```

```
large_values.count()
```

```
9
```

```
large_values.reduce(lambda x,y : x+y)
```

```
135
```

# WordCount in Spark (1 of 2)

```
constitution = sc.textFile("/user/root/constitution.txt")

wordCounts = constitution.flatMap(lambda line: line.split())
                        .map(lambda word: (word, 1))
                        .reduceByKey(lambda a, b: a+b)
```

```
wordCounts.take(10)
```

```
[(u'all', 37),
 (u'Jr.', 1),
 (u'Legislatures', 3),
 (u'Roads;', 1),
 (u'bear', 2),
 (u'needful', 2),
 (u'Place', 3),
 (u'four', 2),
 (u'race,', 1),
 (u'Department', 1)]
```

## WordCount in Spark (2 of 2)

```
swapped_wordCounts = wordCounts.map(lambda record : (record[1], record[0]))
sorted_counts = swapped_wordCounts.sortByKey(ascending = False)
sorted_counts.take(20)
```

```
[(662, u'the'),
 (493, u'of'),
 (293, u'shall'),
 (256, u'and'),
 (183, u'to'),
 (178, u'be'),
 (157, u'or'),
 (137, u'in'),
 (100, u'by'),
 (94, u'a'),
 (85, u'United'),
 (81, u'for'),
 (79, u'any'),
 (72, u'President'),
 (64, u'The'),
 (64, u'as'),
 (63, u'have'),
 (55, u'States,'),
 (52, u'such'),
 (47, u'State')]
```



# Operations on key-value pair RDDs

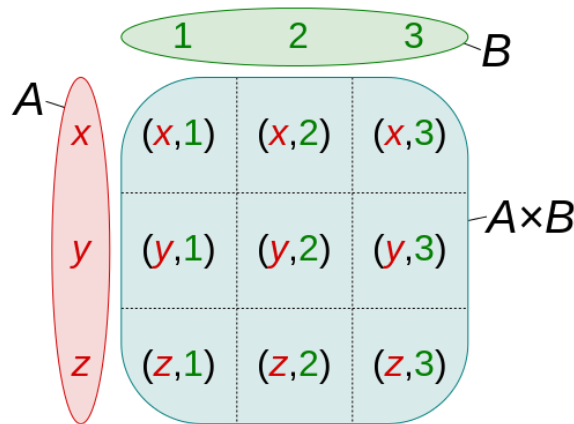
- Operations on key-value pair datasets are at the foundation of Hadoop, MapReduce, and Spark 1.6
  - **groupByKey()**: group values with the same key, ex. `rdd.groupByKey()`
  - **mapValues(f)**: applies function only to values not keys
  - **flatMapValues(f)**: same as `mapValues` when function returns several values
  - **keys()**: Returns RDD with only the keys
  - **values()**: Returns RDD with only the values
- Simple operations on RDDs:
  - **union(otherRDD)**: makes the union of all key-value pairs

# Join operations on pairs of key-value pair RDDs: join operations

- Joins are a fundamental operation of data which compute the **Cartesian product** between two sets (e.g., two RDDs)

$$A \times B = \{(a, b) \mid a \in A \text{ and } b \in B\}$$

- Most of the time, joins are paired with a filter to improve performance



# Join operations on pairs of key-value pair RDDs: join operations

- We can use this idea to *join* key-value pairs and then filter for pairs that have the same key
- RDD join performs an “inner join”
- Inner Join: Keys must be present in the left and right hand RDD.

```
A = sc.parallelize([
    [1,1],
    [1,2],
    [2,3]]
)
B = sc.parallelize([
    [1, "A"],
    [2, "B"]
])
```

```
A.join(B).collect()
```

```
[(1, (1, 'A')), (1, (2, 'A')), (2, (3, 'B'))]
```

# Join operations on pairs of key-value pair RDDs

- Lets suppose we want to compute the total number of orders per state using:
  - Locations: locationID, state
  - Transactions: transactionID, locationID, number of orders
- Activity: how can we use join to achieve this?
- Explore leftOuterJoin and rightOuterJoin

```
locations = sc.parallelize([
    ['loc1', 'NY'],
    ['loc2', 'NY'],
    ['loc3', 'PA'],
    ['loc4', 'FL']
])
transactions = sc.parallelize([
    [1, 'loc1', 2.],
    [2, 'loc1', 3.],
    [3, 'loc2', 5.],
    [4, 'loc5', 5.]
])
```

# Spark 2.0

DataFrames

# DataFrames

- The problem with RDDs is that they do not have enough structure
- They are harder to optimize and therefore slow
- DataFrames tries to solve this by adding structure
- A DataFrame is a distributed collection of data organized into named columns
- Similar to Pandas DataFrames but distributed across the cluster
- You access the Spark 2.0+ functionality using the **spark** object

## DataFrames (2)

- You can read from multiple sources into dataframes



{ JSON }

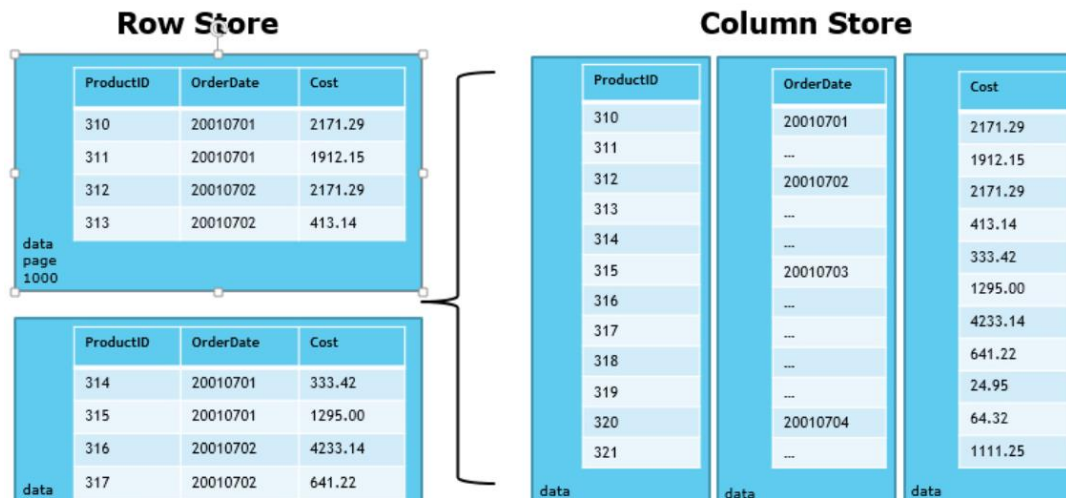


and more ...

Formats and Sources supported by DataFrames

# DataFrames (3)

- One preferred source is Parquet files
- Parquet files are datasets stored in columns





# DataFrame operations

- Creation of dataframes:

- From Row objects:

```
1 from pyspark.sql import Row
2 raw_data = [Row(state='NY', month='JAN', orders=3),
3             Row(state='NJ', month='JAN', orders=4),
4             Row(state='NY', month='FEB', orders=5),
5             ]
6 data_df = spark.createDataFrame(raw_data)
```

- From RDDs:

```
1 example_rdd_with_rows.toDF()
```

```
1 farmers_markets = spark.read.csv('/databricks-datasets/data.gov/farmers_markets_geographic_data/data-001/market_data.csv',
2                                  header=True,
3                                  inferSchema=True)
```

- From files:

```
1 spark.read.parquet(filepath)
```

# DataFrame operations

- All data in a specific column has the same type
- DataFrame types can be hierarchical: The type of a column might be another “DataFrame”
- **You cannot perform dataframe operations using Python**
- Instead, you transform DataFrames by **selecting, modifying, filtering, joining, grouping**, or **aggregating** using specialized Spark commands similar to SQL
- Many of these commands are **symbolic representations** of the operations
- After the transformations, Spark builds an *execution plan* that is optimized for the column types and the operations

# Exploring a DataFrame

- **printSchema()**: shows the datatypes of the dataframe
- **show()**: prints the first  $n$  rows
- **take()**: return first  $n$  rows
- **sample(withReplacement, fraction)**: randomly sample rows (approximate)
- **display(df)**: (only in DataBricks) exploratory interface for dataframe

# Selecting and modifying a DataFrame

- **select(\*expressions)**: returns a new DataFrame with columns
- **withColumn(colName, expression)**: creates a new column based on the expression
- Expressions are **symbolic operations**
- Symbolic operations can hold *literal* values and *placeholders* for column names
- You can perform symbolic operations on both literals and placeholders
- Some symbolic operations are available in `pyspark.sql.functions`

# Symbolic operations

- expression =  $1 + n\_employees$
- To express the previous operation, I need a *literal* value (1) and a *placeholder* for the column `n_employees`

```
1 + fn.col('n_employees')
```

```
t[23]: Column<(n_employees + 1)>
```

- A symbolic operation produces a *column object* which contains the *defined operations*

# Symbolic operations

```
1 locations_df.select(1 + fn.col('n_employees')).show(10)
```

► (3) Spark Jobs

```
+-----+
| (n_employees + 1) |
+-----+
|                4 |
|                9 |
|                4 |
|                2 |
+-----+
```

## Change column name

```
1 new_column = 1 + fn.col('n_employees')
2 locations_df.select(new_column.alias('new_column')).show(10)
```

► (3) Spark Jobs

```
+-----+
|new_column|
+-----+
|        4 |
|        9 |
|        4 |
|        2 |
+-----+
```

## Change column type

```
1 locations_df.\
2   select(new_column.alias('new_column').cast('float')).\
3   show(10)
```

► (3) Spark Jobs

```
+-----+
|new_column|
+-----+
|        4.0 |
|        9.0 |
|        4.0 |
|        2.0 |
+-----+
```

# Selecting and modifying (1)

- You can use select to select columns, modify them, or create new columns

```
1 locations_df.show(10)
```

► (3) Spark Jobs

location_id	n_employees	state
loc1	3	NY
loc2	8	NY
loc3	3	PA
loc4	1	FL

```
1 locations_df.\n2   select('n_employees',\n3         'location_id',\n4         'state',\n5         (fn.col('n_employees') + 1).alias('n_employees_plus_1'),\n6         (fn.col('n_employees') > 5).alias('more_than_5_empl')\n7       ).\n8   show(10)
```

► (3) Spark Jobs

n_employees	location_id	state	n_employees_plus_1	more_than_5_empl
3	loc1	NY	4	false
8	loc2	NY	9	true
3	loc3	PA	4	false
1	loc4	FL	2	false

## Selecting and modifying (2)

- The following code snippet in that new columns are created using utility functions provided by the Spark SQL utility function (fn) module.
- Note that `fn.lit()` is needed to create a literal passed to `pow`

```
1 locations_df.\n2   select(fn.sqrt('n_employees'),\n3         'n_employees',\n4         fn.pow(fn.col('n_employees'),\n5               fn.lit(2))).\n6   show()
```

► (3) Spark Jobs

SQRT(n_employees)	n_employees	POWER(n_employees, 2)
1.7320508075688772	3	9.0
2.8284271247461903	8	64.0
1.7320508075688772	3	9.0
1.0	1	1.0



# Filtering

- **where(expression)**: select only rows where expression is true
- Expressions can be complex:

```
1 locations_df.where((fn.col('n_employees') > 2) & (fn.col('state') == 'PA')).show()
```

► (3) Spark Jobs

location_id	n_employees	state
loc3	3	PA

# Joining

- **Inner joins** (keep rows with keys that exist in the left and right datasets)
- **Outer joins** (keep rows with keys in either the left or right datasets)
- **Left outer joins** (keep rows with keys in the left dataset)
- **Right outer joins** (keep rows with keys in the right dataset)
- **Left semi joins** (keep the rows in the left, and only the left, dataset where the key appears in the right dataset)
- **Left anti joins** (keep the rows in the left, and only the left, dataset where they do not appear in the right dataset)
- **Natural joins** (perform a join by implicitly matching the columns between the two datasets with the same names)
- **Cross (or Cartesian) joins** (match every row in the left dataset with every row in the right dataset)

# Inner Join

- Keep rows with keys that exist in left and right data frames

```
1 locations_df.join(transactions_df, on='location_id').show()
```

► (5) Spark Jobs

location_id	n_employees	state	n_orders	transaction_id
loc1	3	NY	2.0	1
loc1	3	NY	3.0	2
loc3	3	PA	5.0	3

```
1 locations_df.show(10)
```

► (3) Spark Jobs

location_id	n_employees	state
loc1	3	NY
loc2	8	NY
loc3	3	PA
loc4	1	FL

```
1 transactions_df.show(10)
```

► (3) Spark Jobs

location_id	n_orders	transaction_id
loc1	2.0	1
loc1	3.0	2
loc3	5.0	3
loc5	5.0	4

# Left outer join

- Keep rows with keys that exist in left data frame

```
1 locations_df.join(transactions_df, on='location_id', how='left').show()
```

► (5) Spark Jobs

location_id	n_employees	state	n_orders	transaction_id
loc1	3	NY	2.0	1
loc1	3	NY	3.0	2
loc3	3	PA	5.0	3
loc2	8	NY	null	null
loc4	1	FL	null	null

```
1 locations_df.show(10)
```

► (3) Spark Jobs

location_id	n_employees	state
loc1	3	NY
loc2	8	NY
loc3	3	PA
loc4	1	FL

```
1 transactions_df.show(10)
```

► (3) Spark Jobs

location_id	n_orders	transaction_id
loc1	2.0	1
loc1	3.0	2
loc3	5.0	3
loc5	5.0	4

# Right outer join

- Keep rows with keys that exist in right data frame

```
1 locations_df.join(transactions_df, on='location_id', how='right').show()
```

► (5) Spark Jobs

location_id	n_employees	state	n_orders	transaction_id
loc1	3	NY	2.0	1
loc1	3	NY	3.0	2
loc3	3	PA	5.0	3
loc5	null	null	5.0	4

```
1 locations_df.show(10)
```

► (3) Spark Jobs

location_id	n_employees	state
loc1	3	NY
loc2	8	NY
loc3	3	PA
loc4	1	FL

```
1 transactions_df.show(10)
```

► (3) Spark Jobs

location_id	n_orders	transaction_id
loc1	2.0	1
loc1	3.0	2
loc3	5.0	3
loc5	5.0	4

# Outer join

- Keep rows with keys that exist in either the left or right data frames

```
1 locations_df.join(transactions_df, on='location_id', how='outer').show()
```

► (5) Spark Jobs

location_id	n_employees	state	n_orders	transaction_id
loc1	3	NY	2.0	1
loc1	3	NY	3.0	2
loc3	3	PA	5.0	3
loc2	8	NY	null	null
loc4	1	FL	null	null
loc5	null	null	5.0	4

```
1 locations_df.show(10)
```

► (3) Spark Jobs

location_id	n_employees	state
loc1	3	NY
loc2	8	NY
loc3	3	PA
loc4	1	FL

```
1 transactions_df.show(10)
```

► (3) Spark Jobs

location_id	n_orders	transaction_id
loc1	2.0	1
loc1	3.0	2
loc3	5.0	3
loc5	5.0	4

# Grouping

- **groupBy(\*expressions)**: groups by a list of expressions
- Typically, a list of columns:

```
1 locations_df.\n2   join(transactions_df, on='location_id').\n3   groupBy('state')
```

Out[76]: <pyspark.sql.group.GroupedData at 0x7f4ac607ee50>

- Note that this transformation not do anything until we perform an action on the group object like an aggregate action.

# Aggregate

- There are some special functions that only work on grouped data
- They are applied using the method **agg(\*expressions)**
- For example:
  - fn.sum, fn.stddev: self explanatory
  - fn.count: counts when column is not null
  - fn.countDistinct: how many distinct values of a column



# Aggregate

```
1 locations_df.\n2   join(transactions_df, on = 'location_id').\n3   groupBy('state').\n4   agg(fn.sum('n_orders')).\n5   show()
```

## ► (5) Spark Jobs

state	sum(n_orders)
PA	5.0
NY	5.0

# Spark ML

- Spark ML implements several machine learning algorithms at scale:
  - Regression
  - Classification
  - Decision Trees
  - Clustering
- It works on Spark DataFrames by performing transformations of the data
- A typical data science analysis requires several transformations
- These transformations can be implemented through Pipelines

## Spark ML (2)

- A model is known as **Estimator** in Spark ML and the typical cycle for such objects is as follows
  1. Define zero or more **input columns**
  2. Define zero or more **output columns**
  3. Define **parameters** of the estimator
  4. **Fit** the estimator, which returns a **fitted model**
  5. Use the fitted model to perform **transformations**

# Example Spark machine learning workflow

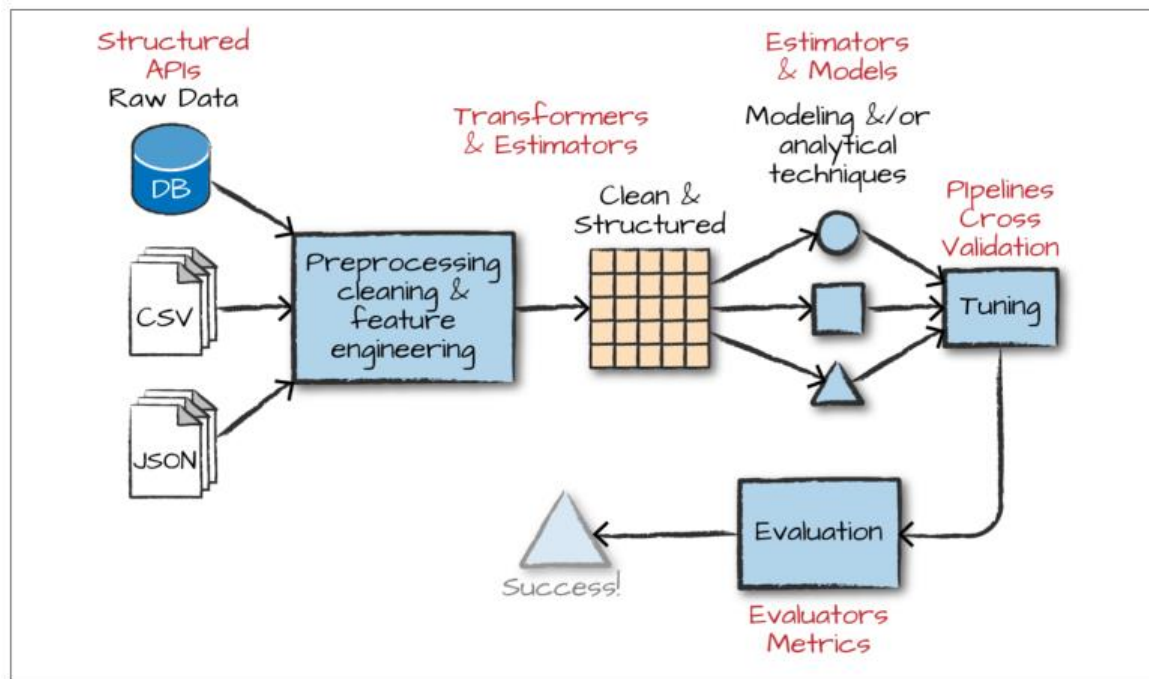
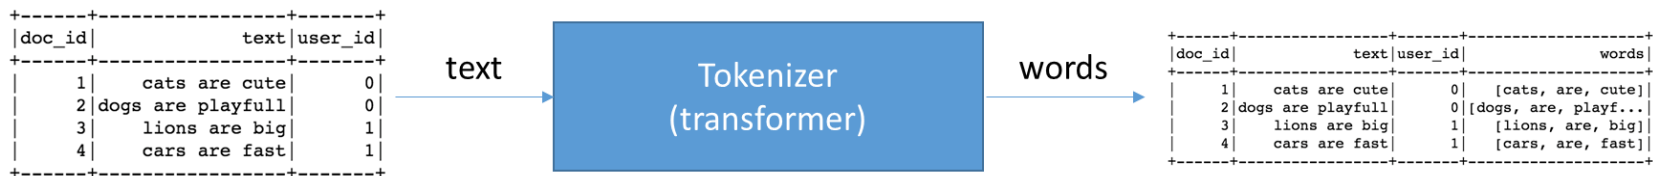


Figure 24-2. The machine learning workflow, in Spark

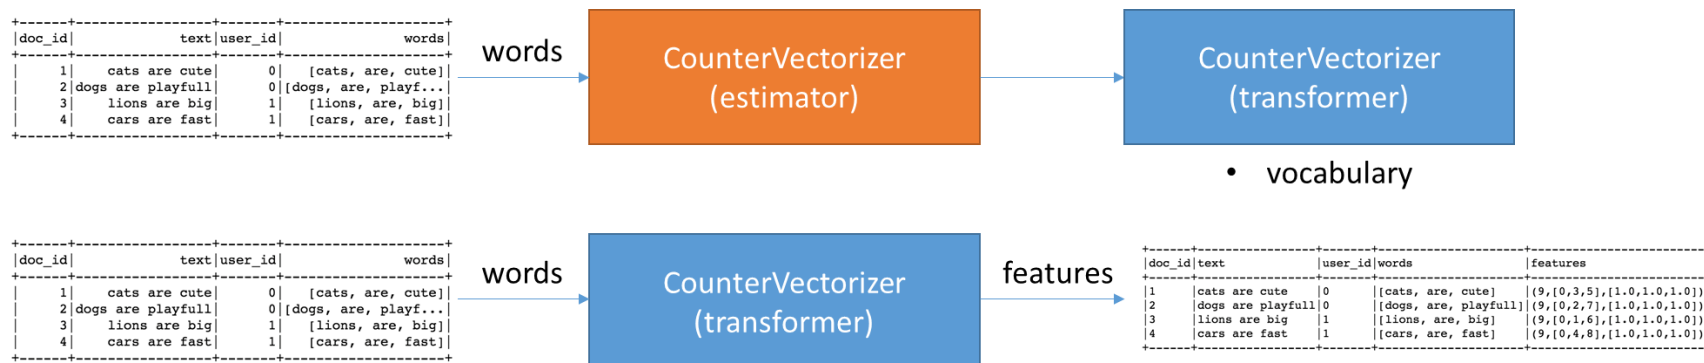
# Spark ML: Transformers and Estimators

- Note that Spark ML for RDD is getting deprecated
- Preferred method in the future is to work only with DataFrames.
- Spark Machine Learning works by creating transformers and estimators DataFrames as inputs

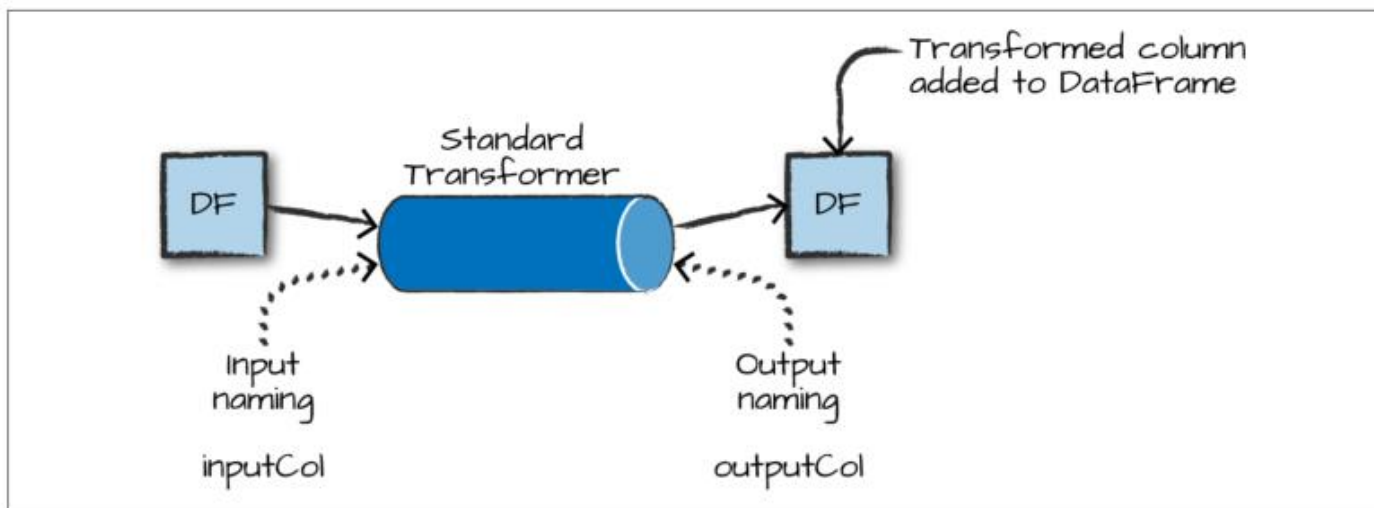


# Spark ML: Transformers and Estimators

- **Transformers** convert raw data in some way
- **Estimators** need to learn something from the data in order to transform the data
- For example, if we want to count terms in text, we need to know how many terms are in all documents



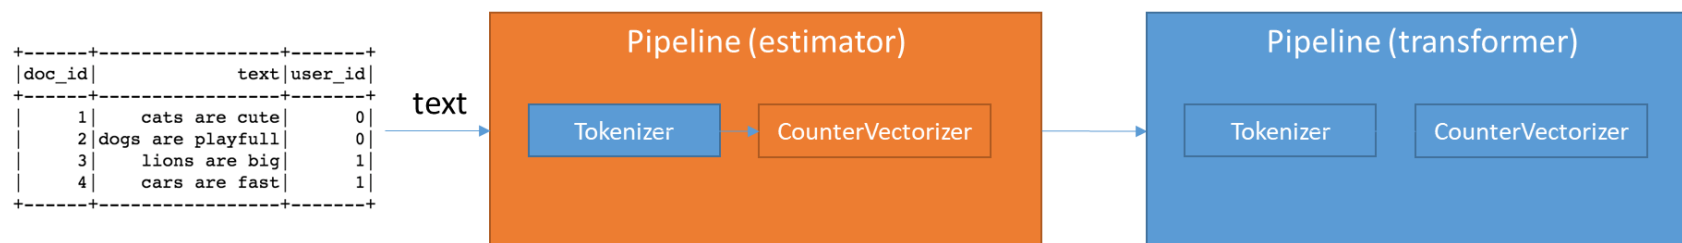
# Example Spark Transformer



*Figure 24-3. A standard transformer*

# Spark ML: Pipelines

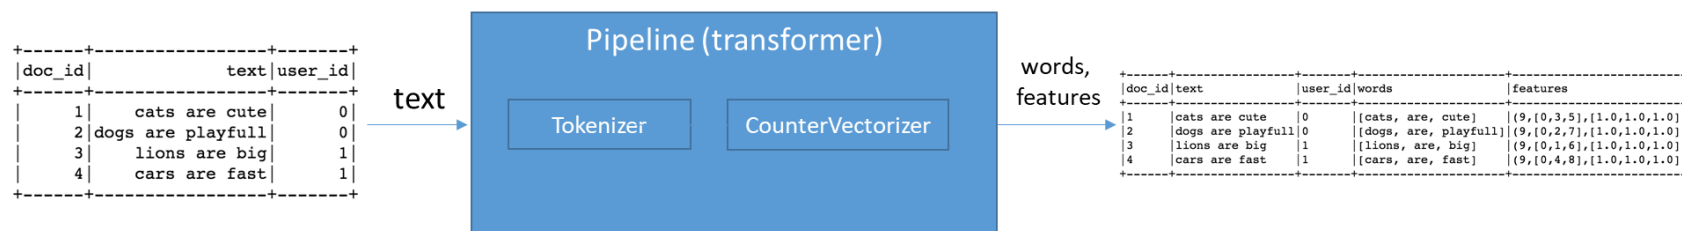
- Data science is all about building data analytic pipelines from raw data to models
- Spark ML Pipeline chains multiple Transformers and Estimators
- Pipelines can be saved and shared
- A Pipeline always starts as an Estimator that needs to be fitted





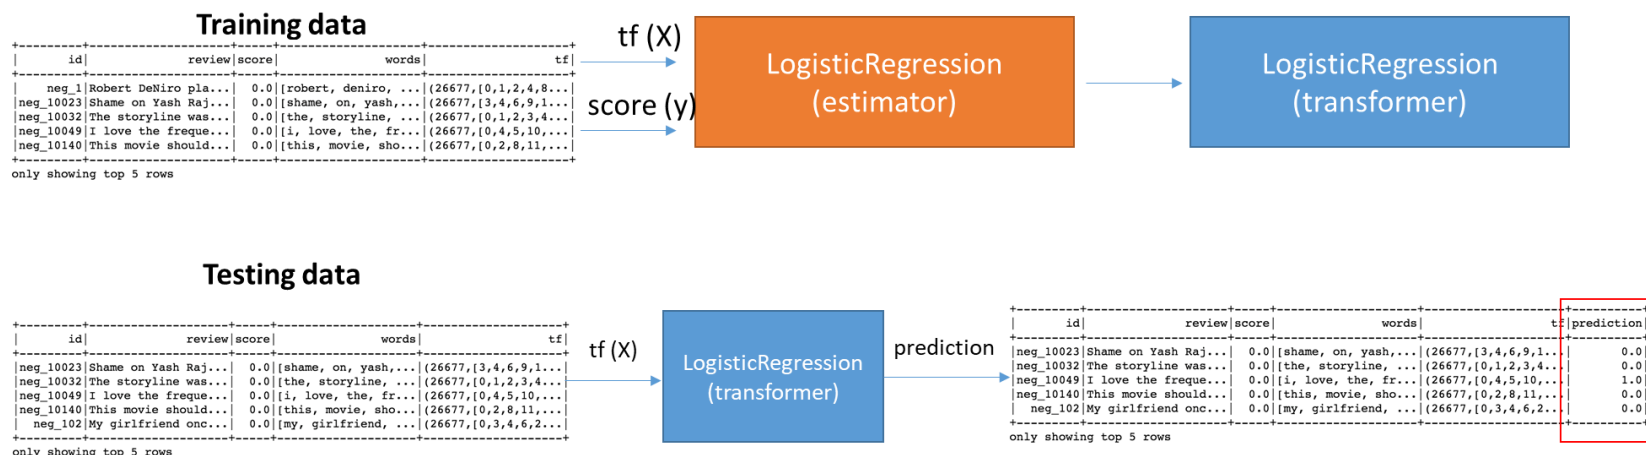
# Spark ML: Pipelines

- A Pipeline **always** needs to be fitted even when stages are all transformers
- A Pipeline transformer contains an entire process



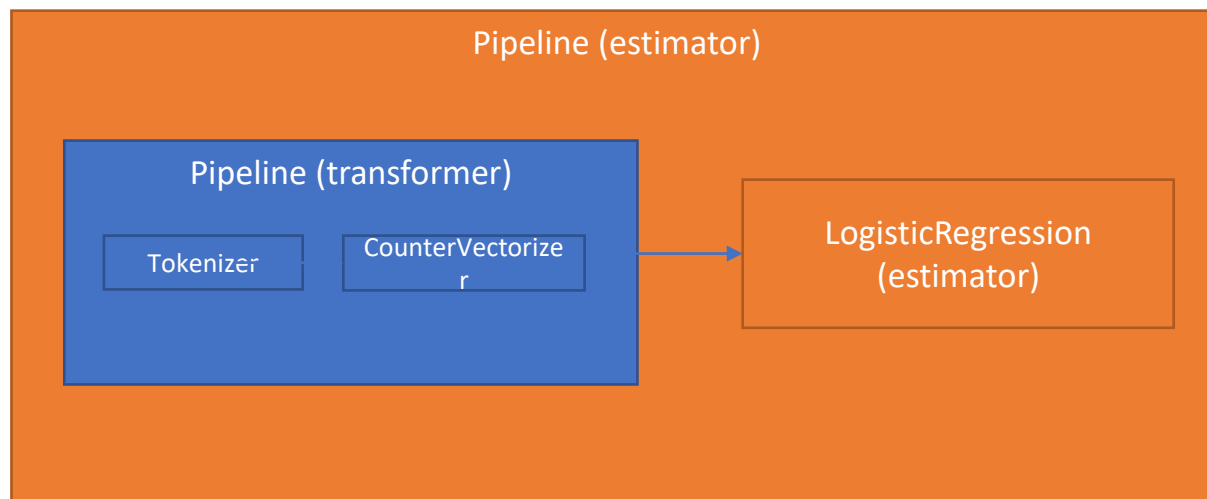
# Spark ML: Algorithms

- Some Estimators are algorithms that work on DataFrames
- In this context, an algorithm is a machine learning model
- For example, LogisticRegression



# Spark ML: Algorithms

- Algorithms can be part of Pipelines
- Pipelines can be combined with other Pipelines



# Evaluators

- An evaluator allows us to see how a given model performs according to specified criteria
- An evaluator is used to select the best model from a group of models.
- An evaluator compares model predictions to truth data and calculates a score indicating how well the model performed.

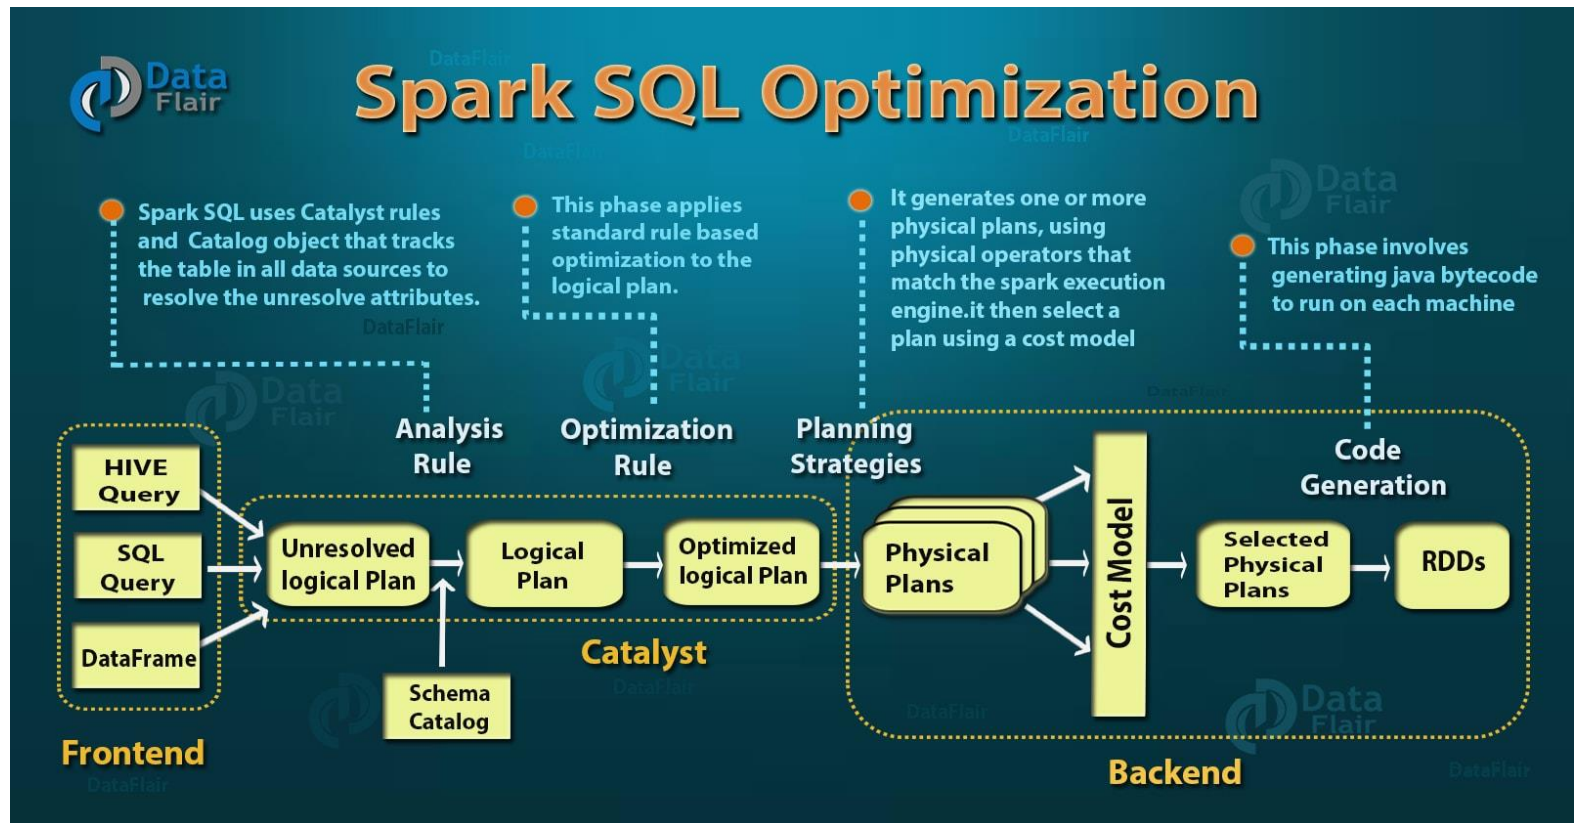
# Summary (1 of 2)

- **Apache Spark** is a fast, in-memory computing system that runs on Hadoop.
- An application in Spark has several main components, including a driver program, worker nodes, executors and tasks.
- The **SparkContext** object is a key component of the driver program.
- An RDD is a fault-tolerant collection of elements that can be operated on in parallel.
- There are several ways to create an RDD in Spark: the **parallelize** function or a reference to a file in HDFS. **textFile sequenceFile**
- There are two types of operations that can be done on RDDs: **transformations** and **actions**.

## Summary (2 of 2)

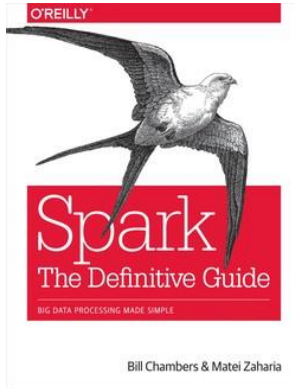
- ***MLlib*** is a Spark implementation of some common machine learning algorithms and utilities.
- MLlib contains **algorithms** for classification, regression, clustering and recommendation.
- **Transformers** convert raw data in some way
- **Estimators** need to learn something from the data in order to transform the data
- **Pipelines** are a collection of transformers and estimators
- Algorithms are

# Spark Sql



# References

---



<https://www.oreilly.com/library/view/spark-the-definitive/9781491912201/>

<https://databricks.com/blog/2016/07/14/a-tale-of-three-apache-spark-apis-rdds-dataframes-and-datasets.html>

<https://data-flair.training/blogs/apache-spark-rdd-vs-dataframe-vs-dataset/>

<http://www.jamesserra.com/archive/2017/12/is-the-traditional-data-warehouse-dead/>